

Error Handling

Ansible Advanced

Agenda

- Error Handling
- Using blocks
- Handlers and Failure
- Execution strategy

Error Handling In Practice

Error Handling

How Ansible behaves with failure

- If you get an error on a task when running on a host, Ansible's default behavior is to continue executing the task on the other hosts
- The task will not break for everyone just because it failed on one host
- However, the playbook execution for the failed host will be halted at the point of the error, and Ansible will not execute subsequent tasks on that specific host
- The other hosts will continue to execute the remaining tasks in the playbook.

Error Handling: ignore_errors

- This directive allows a task to fail without stopping the entire playbook
- It's useful when you expect a task might fail and it's okay if it does.

```
- name: Try to install a package that might not exist
  apt:
    name: some_unknown_package
    state: present
  ignore_errors: yes
```

Error Handling: failed_when

- Allows you to define custom conditions for what constitutes a failure.
- It can be used to override the default failure conditions of a module.it does.

```
- name: Check free disk space
  ansible.builtin.stat:
    path: /opt/data
  register: result
  failed_when: result.stat.size > 5000000
```

Error Handling: any_errors_fatal

- If set to **yes**, any error will be considered fatal and will stop all hosts' execution
- Useful when you want to ensure that if any host fails, no further tasks should be executed on any host
- Can be used on tasks or playbook level

```
- hosts: web_servers
  any_errors_fatal: true
  tasks:
    - name: Ensure configuration file is present
      ansible.builtin.copy:
        src: /src/config.conf
        dest: /dest/config.conf
```

Error Handling: max_fail_percentage

- If you're running a task on multiple hosts, you can use this directive to continue execution even if a certain percentage of hosts fail
- Can be set on task or playbook level

```
- hosts: web_servers
  max_fail_percentage: 30
  tasks:
    - name: Update application
      ansible.builtin.git:
        repo: https://github.com/example/app.git
        dest: /opt/app
```


Unreachable hosts

- Unreachable hosts have a different behavior from failed tasks
- If Ansible cannot connect to a host, it marks that host as 'UNREACHABLE' and removes it from the list of active hosts for the run
- Using **ignore_errors** don't have any effect on unreachable hosts

Use ignore_unreachable

- When you set **ignore_unreachable: yes** for a task or playbook, Ansible will not fail the playbook for that specific task/block if a host becomes unreachable.
- Instead, it will ignore the unreachable host and continue executing the remaining tasks for that host, as well as continue executing tasks for other hosts.
- This can be useful in scenarios where temporary network disruptions are expected, and you don't want the entire playbook run to be affected.

```
- name: Sample Playbook with ignore_unreachable
hosts: all
tasks:
  - name: This task will ignore unreachable hosts
    command: echo "Running a command"
    ignore_unreachable: yes

  - name: This task will fail if a host is unreachable
    command: echo "Another command"
```

Blocks

Error Handling

What are Ansible blocks

- Blocks in Ansible provide a way to group related tasks together
- You may apply particular task attributes, such as looping, conditional execution, and error handling, to the entire group
- They can significantly simplify playbook logic and make error handling more streamlined

Error handling with blocks

- One of the most common uses of blocks is for error handling
- Blocks can be combined with **rescue** and **always** sections to define tasks that should be run if there's an error in the block or tasks that should always run regardless of success or failure

Error handling with blocks

- One of the most common uses of blocks is for error handling
- Blocks can be combined with **rescue** and **always** sections to define tasks that should be run if there's an error in the block or tasks that should always run regardless of success or failure

Basic Error Handling

```
- hosts: localhost
  tasks:
    - block:
        - name: This will fail
          command: /nonexistent/command

    rescue:
      - name: This will run if there's a failure in the block
        debug:
          msg: "An error occurred!"
```

Using **always** for Cleanup Tasks

```
- hosts: localhost
  tasks:
    - block:
        - name: Create a temporary file
          tempfile:
            state: file
            register: temp_file

        - name: This will fail
          command: /nonexistent/command

    rescue:
      - name: Handle the error
        debug:
          msg: "An error occurred!"

    always:
      - name: Cleanup temporary file
        file:
          path: "{{ temp_file.path }}"
          state: absent
```


Combining Blocks with Conditionals and Loops

```
- hosts: localhost
tasks:
  - block:
      - name: Install multiple packages
        apt:
          name: "{{ item }}"
          state: present
        loop:
          - nginx
          - git
        when: ansible_os_family == "Debian"

  rescue:
    - name: Send an alert if installation fails
      debug:
        msg: "Package installation failed!"

  always:
    - name: Always show the OS family
      debug:
        msg: "OS Family: {{ ansible_os_family }}"
```

Demo: Blocks

Ansible Advanced

Handlers and Failure

Error Handling

What are handlers?

- **Purpose:** Handlers are tasks that only run when notified by another task. They are triggered by a "**notify**" directive in a task
- **Idempotency:** Just like regular tasks in Ansible, handlers are idempotent. This means they will only run if the state of the system changes. If no changes are made, the handler won't be triggered
- **Order of Execution:** Handlers run after all tasks are completed in a particular play, not immediately after the task that notified them
- If multiple tasks notify the same handler, the handler will run only once, after all tasks are completed
- The handler only runs if the task that notifies the handler returns a **CHANGED** result

Common use cases

- Restarting a service after its configuration file has been modified
- Reloading a service after a new package has been installed
- Any scenario where a specific action should only be taken in response to a change
- Can be used for error handling

Considerations

- **Handler Naming:** Handlers are identified by their name, so the name specified in the notify directive must match the name of the handler.
- **Order Matters:** If you have multiple handlers, you can control the order in which they run using the **listen** directive. Handlers that listen to the same label will be executed in the order they are defined.
- **Flushing Handlers:** If you want to force all notified handlers to run immediately, rather than waiting until the end of the play, you can use the **meta: flush_handlers** task.

Force handlers on error

- If a task notifies a handler but another task fails later in the play, by default the handler does not run on that host, which may leave the host in an unexpected state
- For example, a task could update a configuration file and notify a handler to restart some service. If a task later in the same play fails, the configuration file might be changed but the service will not be restarted.
- You can change this behavior with
 - **--force-handlers** command-line option
 - Including **force_handlers: True** in a play
 - Adding **force_handlers = True** to **ansible.cfg**
- When handlers are forced, Ansible will run all notified handlers on all hosts, even hosts with failed tasks

Handlers on Error Handling

- Use handlers to perform specific actions in response to errors or failures during playbook execution
- By using blocks, you can group tasks and define error handling for the entire block
- If any task within the block fails, the tasks in the rescue section will execute
- You can notify a handler from within the rescue section to take specific actions in response to errors.

Handlers on Error Handling

```
- hosts: localhost
  tasks:
    - block:
        - name: Update service configuration
          copy:
            src: new_config.conf
            dest: /etc/service/config.conf
          notify: Restart service

    rescue:
        - name: Restore backup configuration
          copy:
            src: backup_config.conf
            dest: /etc/service/config.conf
          notify: Restart service

  handlers:
    - name: Restart service
      service:
        name: myservice
        state: restarted
```

Demo: Handlers

Ansible Advanced

Execution Strategy

Error Handling

Linear Execution

- The linear strategy, which is the default, ensures that each task is executed on all hosts before moving on to the next task.
- If you have a list of tasks and a group of hosts, Ansible will take the first task and execute it on all hosts, one by one (or in parallel batches if you've set a forks value greater than 1).
- Only after the first task is completed on all hosts will Ansible move on to the second task.
- Use Cases
 - When you need to ensure that a particular state is set across all hosts before proceeding.
 - When tasks have dependencies that need to be met on all hosts before moving on.
- Example
 - If you have tasks A, B, and C and hosts 1, 2, and 3, the order of execution will be: A1, A2, A3, B1, B2, B3, C1, C2, C3.

Free Execution

- The free strategy allows each host to process its list of tasks as quickly as possible without waiting for other hosts
- Hosts will execute their list of tasks independently and may finish at different times
- One host might be on task 1 while another host could be on task 3, depending on factors like network speed, host performance, etc.
- Use Cases:
 - When you want to complete the playbook run as quickly as possible and don't have inter-host task dependencies.
 - When working with a large number of hosts where waiting for each task to complete on every host would be inefficient.
- Example
 - If you have tasks A, B, and C and hosts 1, 2, and 3, the order of execution could be: A1, A2, A3, B1, C1, B2, B3, C2, C3. Host 1 moved through all its tasks before Host 2 completed its second task.

Comparison

- Order
 - linear: Ensures a strict order of task execution across all hosts.
 - free: Allows hosts to process tasks as quickly as they can, without a strict order.
- Efficiency
 - linear: Can be slower, especially with a large number of hosts, as each task must complete on all hosts before moving to the next.
 - free: Can be faster as hosts work through tasks independently.
- Predictability
 - linear: More predictable as tasks are executed in a specific order across hosts.
 - free: Less predictable as hosts may be on different tasks at any given time.
- Use Case
 - linear: Best for playbooks where task order and consistency across hosts are crucial.
 - free: Best for playbooks where speed is more important than strict task order.

How to change execution strategy

- Playbook Level

- You can set the strategy for a specific playbook by adding the **strategy** directive at the beginning of your playbook

```
- hosts: all
  strategy: free
  tasks:
```

- Ansible Configuration

- If you want to set the free strategy as the default for all playbooks, you can do so in the Ansible configuration file (ansible.cfg)

```
[defaults]
strategy = free
```

